


METAPHORS OF AI

clustering the metaphors by how dangerous, and how agentic, we imagine the machine to be [27, 45]

- 1Hurricane, tsunami [40]
An overwhelming, impersonal force of nature: no intent, but devastating in its path.
- 2The job-killer [24, 14]
AI as the engine of technological unemployment — the tireless automaton that quietly does the work, and the wages, out of existence.
- 3Slaughterbots [35]
Cheap, autonomous lethal weapons that select and kill targets with no human in the loop.
- 4Pandora's box [6]
A technology that, once opened, releases consequences that cannot be undone or contained.
- 5Ecosystem [30]
AI as a complex, interdependent system that evolves and self-organises beyond any single designer's control.
- 6Wild animal [52]
A powerful, unpredictable, untamed force that can be dangerous if not contained.
- 7A microbiome / organism [23]
AI as a living, growing, distributed organism enmeshed with human society.
- 8Double-edged sword [6]
A tool equally capable of great benefit and great harm, depending on how it is used.
- 9The bullshitter [20]
Fluent text with no regard for truth — indifferent to the facts, not merely mistaken.
- 10Mirror [13, 44]
AI as a reflection of humanity; it shows our own data, biases and culture back to us.
- 11Black box [8, 33]
An opaque system whose inner workings we cannot inspect, only its inputs and outputs.
- 12Child [42]
AI as an immature being we raise and must teach values to; capable, but in need of guidance. Sits near the centre of the agency axis.
- 13The intern [29]
A fast, eager, capable junior that bends the truth — brilliant, never to be trusted unchecked.
- 14Alchemy [34]
Powerful results we cannot explain; a practice that works without a theory of why.
- 15A blurry JPEG of the web [10]
A lossy compression of everything written online; fluent approximation that blurs the facts it restates.
- 16Domesticated animal [39]
A tamed, useful, mostly predictable companion that still keeps its own drives.
- 17The therapist (ELIZA) [47]
A mirror that listens; the confidant we over-trust, projecting understanding onto a simple program.
- 18Tool (hammer) [25, 50]
AI as a neutral instrument; its value and danger come entirely from how a human wields it.
- 19The Mechanical Turk [17]
Seemingly autonomous intelligence quietly powered by an invisible workforce of human labour.
- 20Butler or slave [7, 9]
An obedient servant that performs labour on command; raises questions of autonomy and exploitation.
- 21Parrot [3]
The “stochastic parrot”: AI that mimics and recombines language without genuine understanding.



- 22The bomb [2]
AI as the new strategic weapon — a national-security race run on the logic of the Manhattan Project.
- 23Puppet master [53]
AI covertly pulling the strings, manipulating human behaviour and society from behind the scenes.
- 24Frankenstein's monster [38]
The creature that destroys its creator; the dread of building something we cannot answer for.
- 25Robots ruling society [9]
Embodied machines taking control of social and political structures.
- 26Menacing overlord [5, 16]
A hostile superintelligence that deliberately dominates and rules over humanity.
- 27The sleeper agent [21]
A model that behaves until triggered — deception trained so deep it survives our attempts to remove it.
- 28Human squishing an ant [5, 51]
A superintelligence so far beyond us that it harms humans with indifference, the way a person steps on an ant without malice.
- 29The golem [49]
The clay servant brought to life that may turn on its maker — the oldest cybernetic cautionary tale.
- 30The Singularity [16, 46, 26]
The hypothesised point where machine intelligence outruns our own and triggers runaway, unknowable change — AGI and ASI as an event horizon beyond which prediction breaks down.
- 31Genie [36, 48]
A wish-granting power that delivers what you literally ask for, with the risk of unintended interpretations.
- 32King Midas [36, 48]
AI that gives you exactly what you specify, turning everything to “gold” with catastrophic side effects; the specification / alignment problem.
- 33Brain [4, 37]
AI as a disembodied mind, a thinking organ with potentially superhuman cognition. Sits near the top of the agency axis.
- 34The code monkey [22, 11]
AI as the autonomous software developer — Devin and the wave of agentic coding tools that write, debug and ship code on their own.
- 35The cyborg [19]
The dissolving boundary between human and machine; intelligence as prosthesis and fusion.
- 36The Oracle [5, 41]
AI consulted as a seer — the super-forecaster and world-simulator we ask to foretell the future, trusting its predictions as prophecy.
- 37Co-pilot [28]
AI as an assistant working alongside humans, augmenting rather than replacing; the human stays in control.
- 38The artificial lover [43]
AI as intimate companion — the AI girlfriend, Character.AI and Replika: a system designed to be loved, and to be confided in.
- 39A country of geniuses [1]
A nation of brilliant minds in a datacenter, poised to cure disease and remake the world — if it stays aligned.
- 40The centaur [12]
Human and machine fused into one team that outperforms either alone.
- 41AI as “dog-owner” to humans [51]
Role reversal where AI is the master and humans the well-kept pets, comfortable but subordinate.
- 42The new electricity [31]
A general-purpose utility, like electric power, poised to transform every industry.
- 43Guardian angel [15]
A benevolent, protective intelligence that safeguards humans and intervenes on our behalf.
- 44Supernatural entity / benevolent God [18, 32]
AI imagined as an omniscient, godlike higher power that is fundamentally good and watches over humanity.

REFERENCES & FURTHER READING

The metaphors are drawn from the literature; each entry lists a source and a short link to it.

[1] Amodei, D. (2024). “Machines of Loving Grace: How AI Could Transform the World for the Better.” Essay, [darioamodei.com](#). [darioamodei.com/machines-of-loving-grace](#)

[2] Aschenbrenner, L. (2024). Situational Awareness: The Decade Ahead. [situational-awareness.ai](#)

[3] Bender, E. M., Gebru, T., McMillan-Major, A., & Shmitchell, S. (2021). “On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?” Proc. ACM FAccT ’21, 610–623. [doi.org/10.1145/3442188.3445922](#)

[4] Berkeley, E. C. (1949). Giant Brains, or Machines That Think. New York: John Wiley & Sons. [archive.org](#)

[5] Bostrom, N. (2014). Superintelligence: Paths, Dangers, Strategies. Oxford University Press. [global.oup.com](#)

[6] Brundage, M., et al. (2018). “The Malicious Use of Artificial Intelligence.” [arXiv:1802.07228](#). [arxiv.org/abs/1802.07228](#)

[7] Bryson, J. J. (2010). “Robots Should Be Slaves.” In Close Engagements with Artificial Companions, 63–74. John Benjamins. [benjamins.com](#)

[8] Burrell, J. (2016). “How the Machine ‘Thinks’: Understanding Opacity in Machine Learning Algorithms.” Big Data & Society, 3(1). [doi.org/10.1177/2053951715622512](#)

[9] Čapek, K. (1920). R.U.R. (Rossum's Universal Robots). Prague: Aventinum. [gutenberg.org/ebooks/59112](#)

[10] Chiang, T. (2023). “ChatGPT Is a Blurry JPEG of the Web.” The New Yorker, 9 Feb. [newyorker.com](#)

[11] Cognition AI (2024). “Introducing Devin, the First AI Software Engineer.” [cognition.ai/blog/introducing-devin](#)

[12] Cowen, T. (2013). Average Is Over: Powering America Beyond the Age of the Great Stagnation. New York: Dutton. [penguinrandomhouse.com](#)

[13] Crawford, K. (2021). Atlas of AI. New Haven: Yale University Press. [yalebooks.yale.edu](#)

[14] Frey, C. B., & Osborne, M. A. (2017). “The Future of Employment.” Technological Forecasting & Social Change, 114, 254–280. [doi.org/10.1016/j.techfore.2016.08.019](#)

[15] Goertzel, B. (2012). “Should Humanity Build a Global AI Nanny..?” Journal of Consciousness Studies, 19(1–2), 96–111. [ingentaconnect.com](#)

[16] Good, I. J. (1965). “Speculations Concerning the First Ultrapotent Machine.” Advances in Computers, 6, 31–88. [sciencedirect.com](#)

[17] Gray, M. L., & Suri, S. (2019). Ghost Work: How to Stop Silicon Valley from Building a New Global Underclass. Boston: HMH. [ghostwork.info](#)

[18] Harari, Y. N. (2016). Homo Deus: A Brief History of Tomorrow. London: Harvill Secker. [yhharari.com](#)

[19] Haraway, D. (1985). “A Cyborg Manifesto.” Socialist Review, 80, 65–108. [archive.org](#)

[20] Hicks, M. T., Humphries, J., & Slater, J. (2024). “ChatGPT Is Bullshit.” Ethics and Information Technology, 26(2), 38. [doi.org/10.1007/s10676-024-09775-5](#)

[21] Hubinger, E., et al. (2024). “Sleepers: Training Deceptive LLMs that Persist Through Safety Training.” Anthropic. [arXiv:2401.05566](#). [arxiv.org/abs/2401.05566](#)

[22] Jiménez, C. E., et al. (2024). “SWE-bench: Can Language Models Resolve Real-World GitHub Issues?” ICLR. [arxiv.org/abs/2310.06770](#)

[23] Kelly, K. (2010). What Technology Wants. New York: Viking. [k.k.org](#)

[24] Keynes, J. M. (1930). “Economic Possibilities for Our Grandchildren.” In Essays in Persuasion (1931). [gutenberg.net.au](#)

[25] Kranzberg, M. (1986). “Technology and History: ‘Kranzberg’s Laws.’” Technology and Culture, 27(3), 544–560. [doi.org/10.2307/3105385](#)

[26] Kurzweil, R. (2005). The Singularity Is Near. New York: Viking. [penguinrandomhouse.com](#)

[27] Lakoff, G., & Johnson, M. (1980). Metaphors We Live By. University of Chicago Press. [press.uchicago.edu](#)

[28] Licklider, J. C. R. (1960). “Man-Computer Symbiosis.” IRE Trans. Human Factors in Electronics, HFE-1, 4–11. [ieeexplore.ieee.org](#)

[29] Mollick, E. (2024). Co-Intelligence: Living and Working with AI. New York: Portfolio. [penguinrandomhouse.com](#)

[30] Nardi, B. A., & O’Day, V. L. (1999). Information Ecologies: Using Technology with Heart. MIT Press. [mitpress.mit.edu](#)

[31] Ng, A. (2017). “Artificial Intelligence Is the New Electricity.” Keynote, Stanford GSB. [gsb.stanford.edu](#)

[32] O’Gieblyn, M. (2021). God, Human, Animal, Machine. New York: Doubleday. [penguinrandomhouse.com](#)

[33] Pasquale, F. (2015). The Black Box Society. Cambridge, MA: Harvard University Press. [hup.harvard.edu](#)

[34] Rahimi, A. (2017). “Machine Learning Has Become Alchemy.” NeurIPS Test-of-Time Award lecture. [nips.cc](#)

[35] Russell, S., et al. (2017). Slaughterbots [short film]. Future of Life Institute. [autonomousweapons.org](#)

[36] Russell, S. (2019). Human Compatible: Artificial Intelligence and the Problem of Control. New York: Viking. [humancompatible.ai](#)

[37] Shanahan, M. (2023). “Talking About Large Language Models.” [arXiv:2212.03551](#) (Comm. ACM, 2024). [arxiv.org/abs/2212.03551](#)

[38] Shelley, M. (1818). Frankenstein; or, The Modern Prometheus. London: Lackington et al. [gutenberg.org/ebooks/84](#)

[39] Silverstone, R., & Haddon, L. (1996). “Design and the Domestication of ICTs.” In Communication by Design. Oxford UP. [oup.com](#)

[40] Suleyman, M., with Bhaskar, M. (2023). The Coming Wave. New York: Crown. [the-coming-wave.com](#)

[41] Tetlock, P. E., & Gardner, D. (2015). Superforecasting: The Art and Science of Prediction. New York: Crown. [penguinrandomhouse.com](#)

[42] Turing, A. M. (1950). “Computing Machinery and Intelligence.” Mind, 59(236), 433–460. [doi.org/10.1093/mind/LIX.236.433](#)

[43] Turkle, S. (2011). Alone Together. New York: Basic Books. [sherryturkle.com](#)

[44] Vallor, S. (2024). The AI Mirror. Oxford University Press. [global.oup.com](#)

[45] Vickers, B., & Allado-McDowell, K. (eds.) (2020). Atlas of Anomalous AI. London: Ignota Books. [ignota.org](#)

[46] Vinge, V. (1993). “The Coming Technological Singularity.” VISION-21 Symposium, Science, 131(3410), 1355–1358. [doi.org/10.1126/science.131.3410.1355](#)

[47] Weizenbaum, J. (1976). Computer Power and Human Reason. San Francisco: W. H. Freeman. [archive.org](#)

[48] Wiener, N. (1960). “Some Moral and Technical Consequences of Automation.” Science, 131(3410), 1355–1358. [doi.org/10.1126/science.131.3410.1355](#)

[49] Wiener, N. (1964). God & Golem, Inc. Cambridge, MA: MIT Press. [mitpress.mit.edu](#)

[50] Winner, E. (1980). “Do Artifacts Have Politics?” Daedalus, 109(1), 121–136. [jstor.org](#)

[51] Wozniak, S. (2015). Remarks at the Freescale Technology Forum (reported in Time, June 2015). [time.com](#)

[52] Yampolskiy, R. V. (2020). “On Controllability of Artificial Intelligence.” [arXiv:2008.04071](#). [arxiv.org/abs/2008.04071](#)

[53] Zuboff, S. (2019). The Age of Surveillance Capitalism. New York: PublicAffairs. [publicaffairsbooks.com](#)